Welcome to this Training Session with Theiagen Genomics

We will soon be getting started

Software Development Practices for Public Health Bioinformatics

Week 04: Workflow Managers and Nextflow

A Northeastern Bioinformatics Regional Resource Offering Provided by the Massachusetts Department of Public Health in Collaboration with Theiagen Genomics

Course Introduction

Training Workshop Instructors

Andrew Hale, BSc

- Bioinformatics Developer at Theiagen Genomics since 2025
- BSc in Computational Bioinformatics

Course Agenda

Course Agenda

Use of Workflow Managers

Week 4 - March 17/19, 2025

- Basics of workflow managers, with a focus on Nextflow
  - Package managers
  - Workflow managers
  - Nextflow

Package Managers

CONDA

Conda is a powerful command line tool for **package and environment management** that runs on Windows, macOS, and Linux.

Anaconda - Comes with a lot of pre-installed packages

Miniconda - Comes with the bare necessities

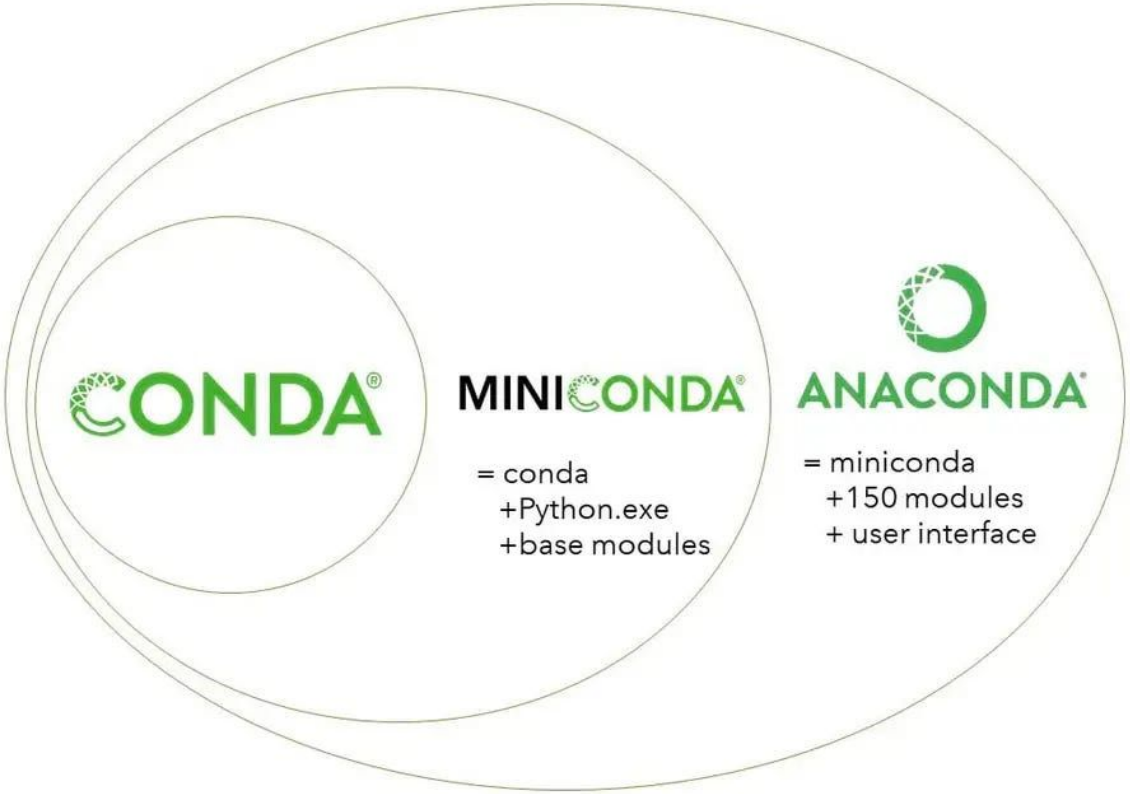
CONDA®
MINICONDA®
= conda
+Python.exe
+base modules
ANACONDA®
= miniconda
+150 modules
+ user interface

CONDA

Why Conda?

- Non-Root access → Use non-root user
- Dependency conflicts → Installs in separate environments
- Shareable → Installs are modular and easy to share
- Easy to execute → Easy to load/unload environments

```
# Download Miniconda installer
wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh

# Run installer
bash Miniconda3-latest-Linux-x86_64.sh
```

Conda Rules of Thumb

There is an exception to this rule of thumb though!

Rule Number 1 - Keep base clean

Keep your base environment clean! Try to avoid installing anything in your base environment. If your base environment breaks, you have to reinstall Miniconda.

Rule Number 2 - Create environments

`conda create` is your friend, use it for everything! Treat environments as consumables, create them, install in them, delete them.

Rule Number 3 - Use containers for critical tasks

Conda is great, but for critical things its best to use containers (Docker or Singularity).

Workflow Managers

Introduction to Workflow Managers

Workflow managers provide a **standardized framework** for creating reproducible and interoperable **bioinformatics pipelines**, especially when **containerized software are utilized**

Workflow managers help to **standardize the process** of writing bioinformatics pipelines

Containerized software will be covered in a future training

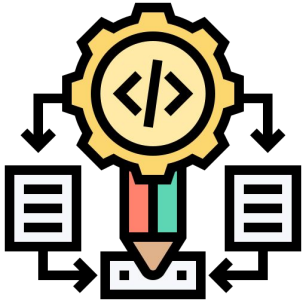

Analytical Module
Individual step performing a single and specific bioinformatics job

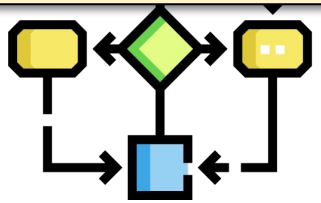

Bioinformatics Pipeline
Software that strings together multiple analytical modules

Theiagen

Introduction to Workflow Managers

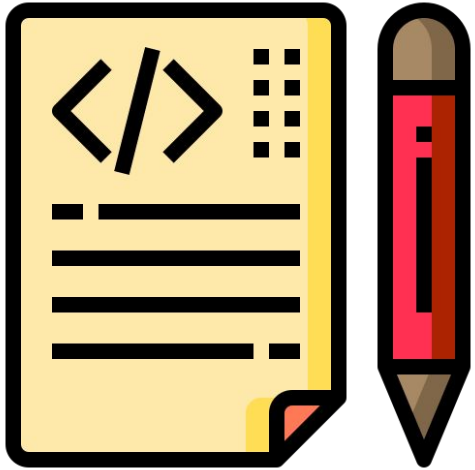

Workflow Language
Programming language used to describe the bioinformatics pipeline

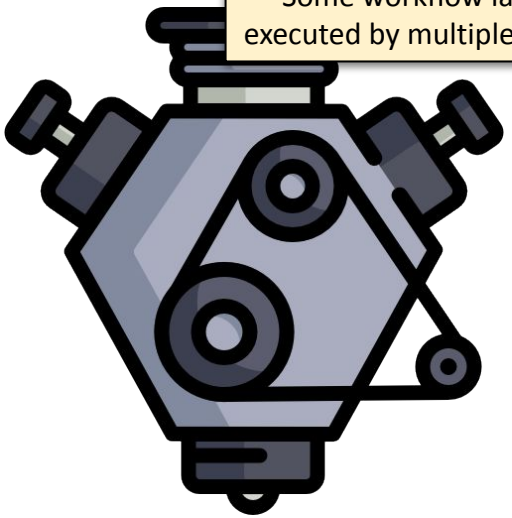

Some workflow languages can be executed by multiple workflow engines

Workflow Engine
Software to interpret and execute the workflow itself

Theiagen

Nextflow

Introduction to Nextflow

Workflows
Using computers to collect, store, analyze and disseminate data and information

Large files
> 10 GB for a metagenomic sample

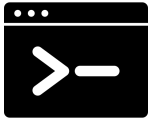

Many languages
Bash, Python, PERL…

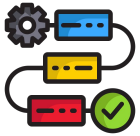

Complex interactions
Network of software and their dependencies

Theiagen

Introduction to Nextflow

Nextflow

Managing modern workflows is complicated

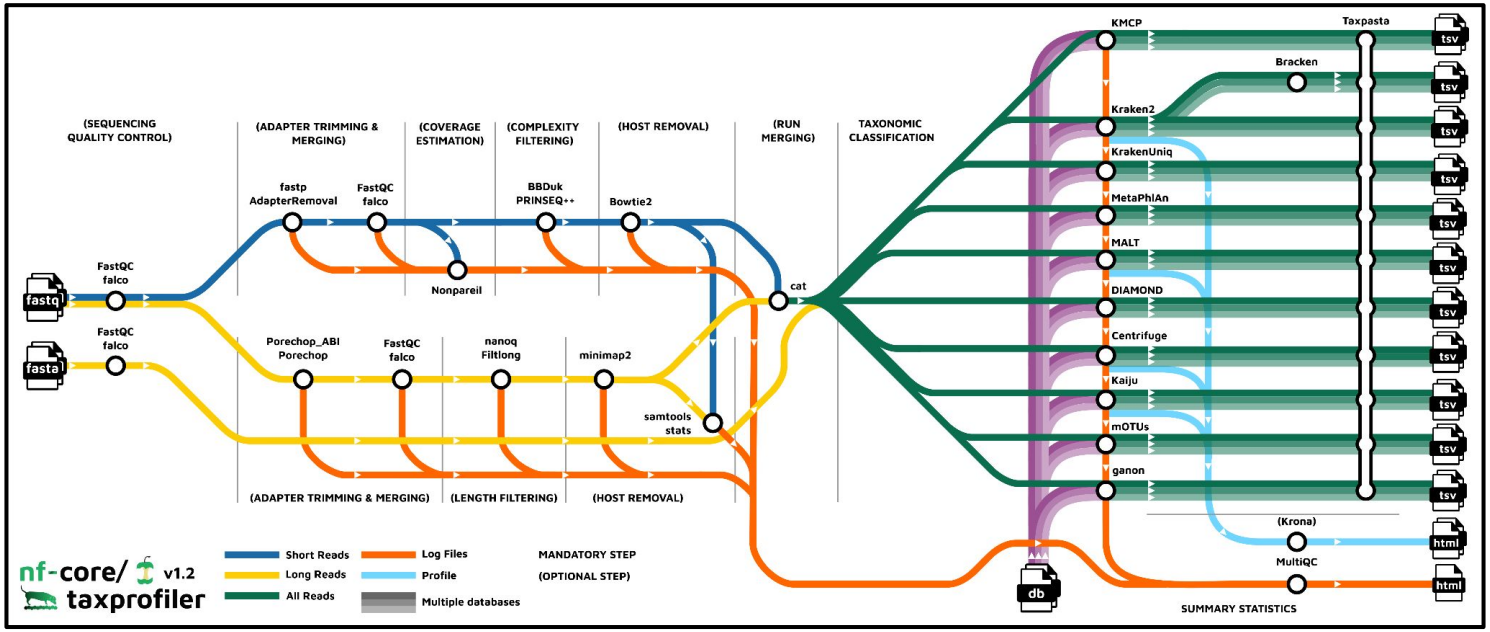

https://nf-co.re/taxprofiler/1.2.2/

Introduction to Nextflow

Reactive workflow framework

Create pipelines with asynchronous data streams

Programing DSL

Has its own language for building a pipeline

Containerized

Out of the box integrating with container engines

Simplifies the deployment of complex parallel and reactive workflows with resumability functionality

```
process sayHello {
  input:
    val cheers
  output:
    stdout

  """
  echo $cheers
  """
}

workflow {
  channel.of('Ciao','Hello','Hola') | sayHello | view
}
```

Theiagen

Introduction to Nextflow

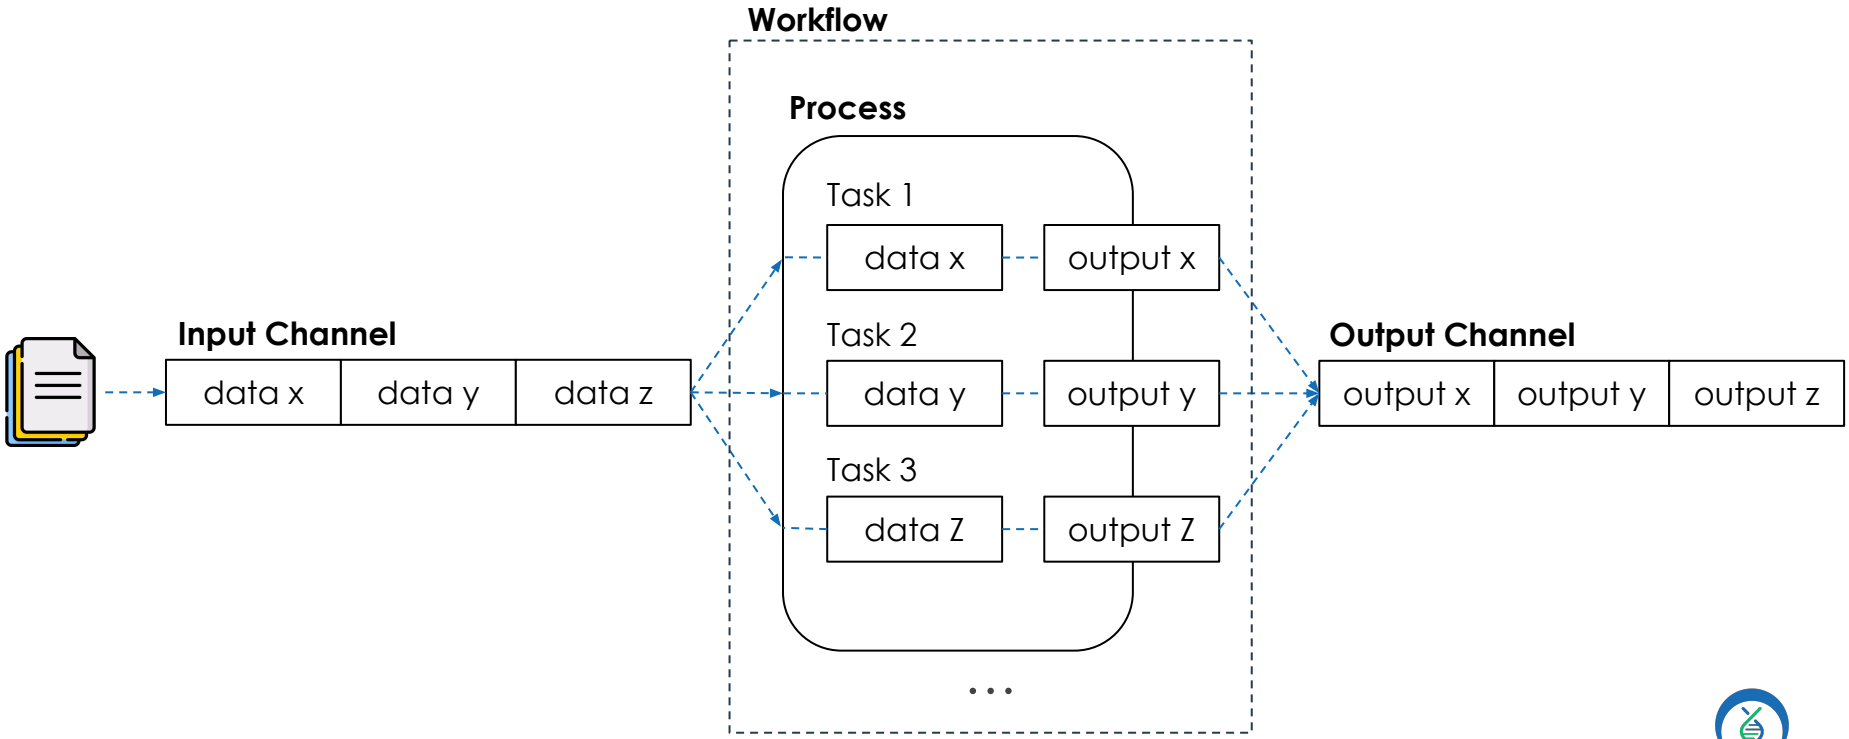

Theiagen

Introduction to Nextflow

Nextflow *Pipeline*

Write code in any language

Orchestrate tasks with dataflow programming

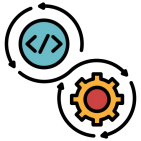

Define software dependencies via containers

Built-in version control with git

Nextflow *Runtime*

Task orchestration and execution

Theiagen

1. Nextflow Basics

Nextflow Basic Structure

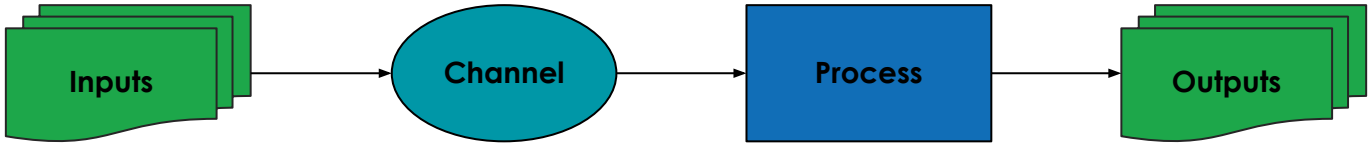

- Channel
  - Queue Channel
    - "First-in, First-out" (FIFO) connecting processes and operators
  - Value Channel
    - Stores a single value (e.g., genome size)
    - Can be a named list (e.g., ["id": "my_sample", "genome_size": 360000])

- Channels can be used by **Operators** and **Processes**

Theiagen

Nextflow Basic Structure

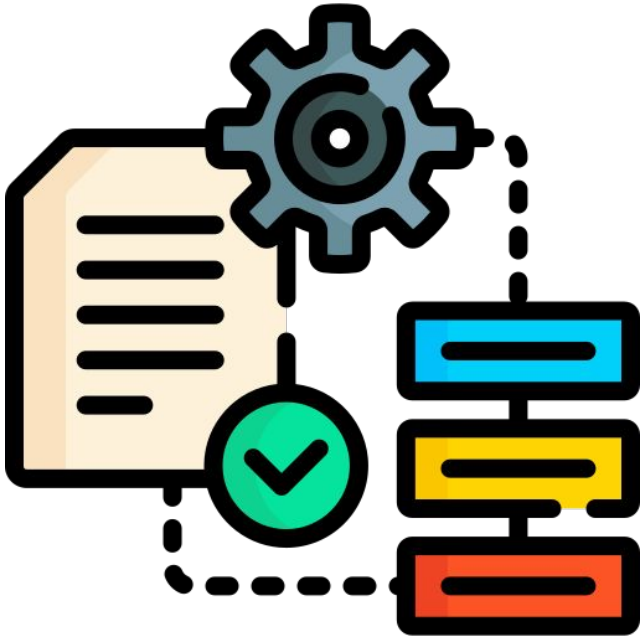

- Methods to **connect** channels, or **transform** values of channels
  - Filtering: filter, randomSample, take, unique
  - Transforming: collect, groupTuple, reduce
  - Text Processing: splitCsv, splitJson, splitText
  - Combining: combine, concat, join, mix
  - Forking: branch, multiMap
  - Math: count, max, min, sum
  - Other: ifEmpty, map, set, view
- More than 50 Operators available to use

Theiagen

Nextflow Basic Structure

Process example:

```
bwa mem reference.fa sample.fq \
    | samtools sort -o sample.bam
```

Theiagen®

Nextflow Basic Structure

Process example:

```
process align_sample {

    input:
    file 'reference.fasta'
    file 'sample.fq'

    output:
    file 'sample.bam'

    script:
    """
    bwa mem reference.fa sample.fq \
        | samtools sort -o sample.bam
    """
}
```

Theiagen

Nextflow Basic Structure

Processes:

- The basic unit for executing user scripts
- 30+ directives adjust optional settings
- Can be dynamic
- Inputs and outputs are channels
- Can be optional
- Conditional executions using ‘when’
- The script block executes user code

```
process < name > {

   [ directives ]

   input:
    < process inputs >

   output:
    < process outputs >

   when:
    < condition >

   [script|shell|exec]:
   < user script to be executed >

}
```

Nextflow Basic Structure

FIFO Can Cause "Unexpected" Results:

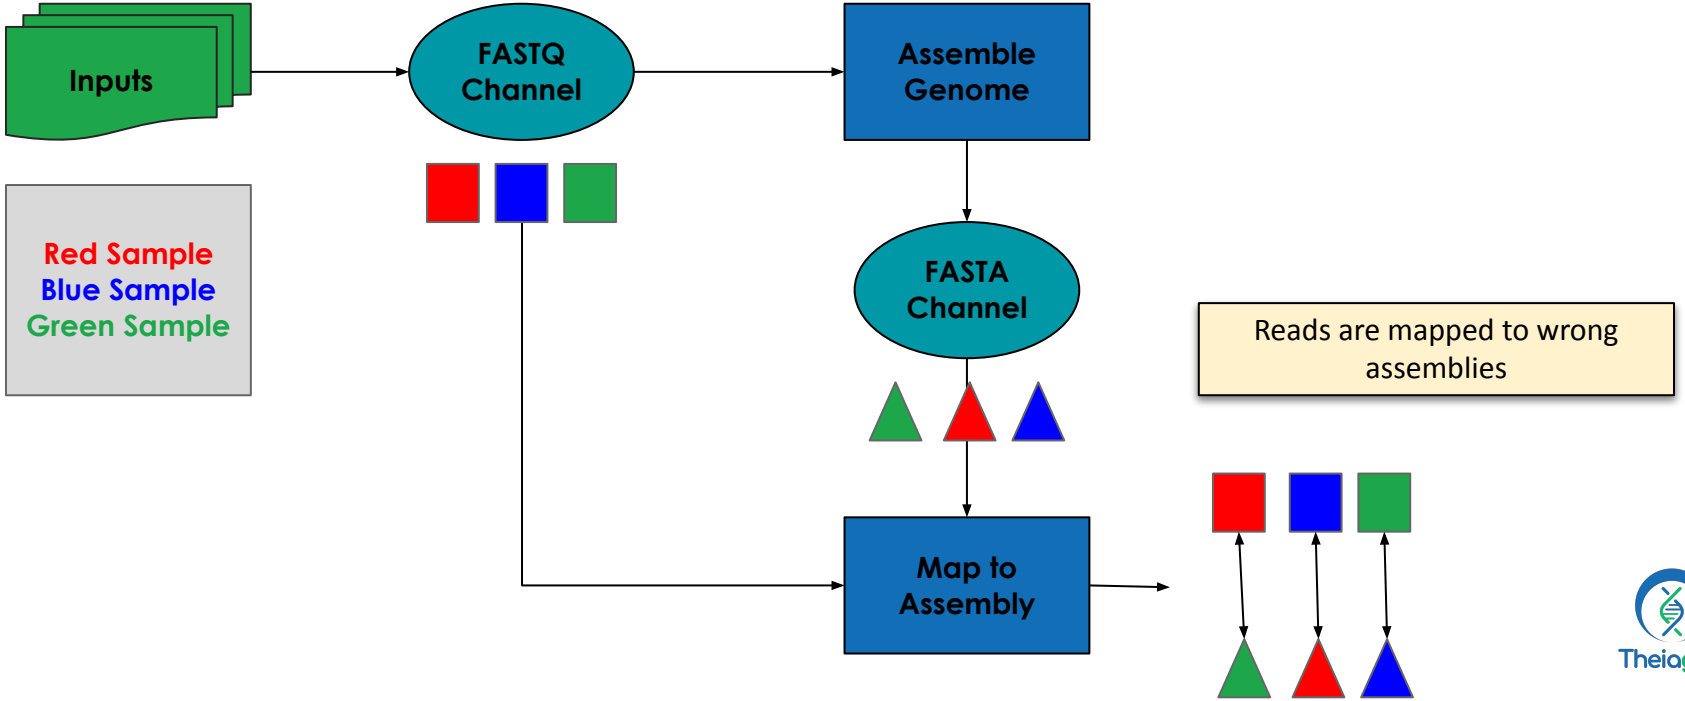

Theiagen

Nextflow Basic Structure

FIFO Can Cause "Unexpected" Results:

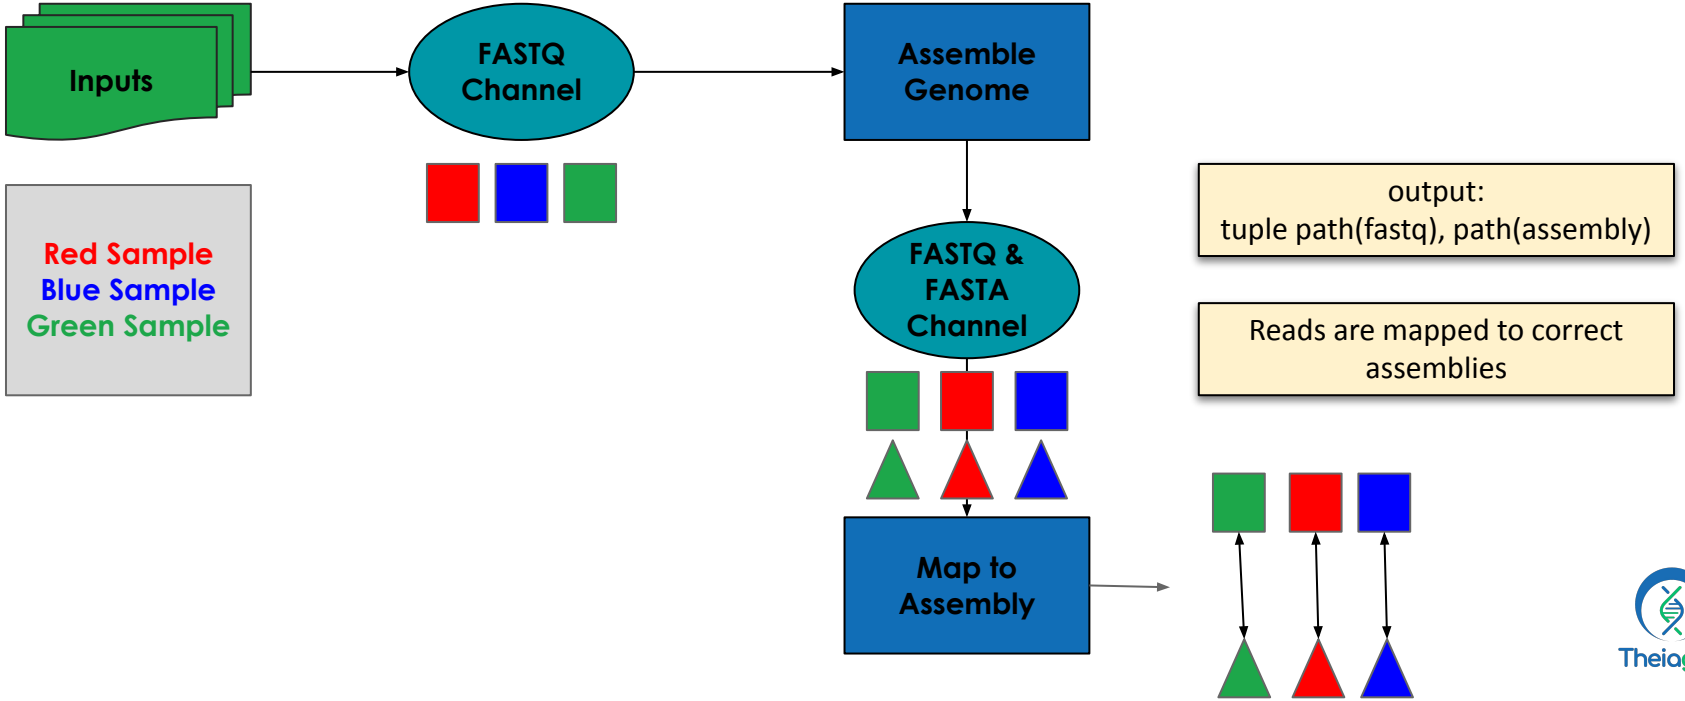

Theiagen

2. Nextflow Files

Nextflow Files

The Nextflow module file

- Modules – A **reusable** Nextflow script with a process definition
- Subworkflows – Multiple modules linked together

Modules are portable and easily shared between workflows

```
process < name > {

   [ directives ]

   input:
    < process inputs >

   output:
    < process outputs >

   when:
    < condition >

   [script|shell|exec]:
   < user script to be executed >

}
```

Theiagen

Nextflow Files

The Nextflow script file
The `main.nf` file is the central script in a Nextflow workflow.

It defines the structure and execution logic of the pipeline by **orchestrating the processes and connecting them with channels**. This file acts as the "blueprint" for how data flows through the pipeline and how tasks are executed.

```
#!/usr/bin/env nextflow
nextflow.enable.dsl=2

// Correctly include the process
// definition from the module
include { newModule } from
'./modules/new_module'

workflow {
    // Define the input channel from the
    // user-specified path
    IN_FilePath =
Channel.fromPath(params.input).ifEmpty {
        exit 1, "No file provided with
pattern: ${params.input}"
    }

    // Execute the 'newModule' process
    newModule(IN_FilePath) | view
}
```

Nextflow Files

The Nextflow config file

When a pipeline script is launched, Nextflow looks for configuration files. By default it searches for `nextflow.config` but one can be provided directly via the `-c <config-file>` optional parameter.

```
// Define default settings
params {
    input = null
}

// Configure process settings
process {
    cpus = 2
    memory = '2 GB'
    time = '1h'
}

profiles {
    standard {
        // Default profile for local execution
        process.executor = 'local'
    }

    docker {
        process {
            executor = 'docker'
            container = 'ubuntu:jammy'
        }
    }
}
```

Nextflow Files

Putting everything together
This is not obligatory, but following this structure is highly recommended.

The modules live in a `modules` folder where the `main.nf` and `nextflow.config` files are.

```
bin/nextflow/
├── main.nf
├── modules
│   └── fastq_stats.nf
└── nextflow.config
```

Theiagen

3. Running Nextflow Workflows

Running Nextflow Workflows

```
nextflow run main.nf -c nextflow.config --input sample.fastq
```

To run a pipeline:

- All software dependencies must be installed (Nextflow + Docker / Singularity / Conda).
- Configure Nextflow to run on your system (-profile)
- Run the tests for your pipeline in the terminal to confirm everything is working (-profile test)

Core Nextflow command-line options use one (-), whereas pipeline-specific parameters use two (--).

Use -resume to restart pipelines that did not complete. This uses cached results for successful tasks from the previous run, instead of executing all tasks from scratch.

5. Nf-Core

nf-core

Develop with a community

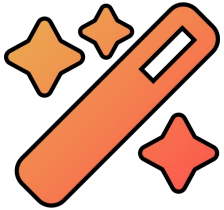

Use a common template

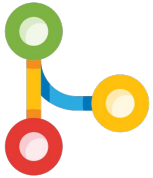

Collaborate, don't duplicate

Nf-Core Github Repository

nf-core

A community effort to collect a curated set of analysis pipelines built using Nextflow.

Verified

Follow

1.7k followers · http://nf-co.re · @nf-co.re · @nf_core@mstdn.science · company/nf-core · core@nf-co.re

README.md

nf-core

Welcome to @nf-core!

@nf-core is a community effort to collect a curated set of analysis pipelines built using Nextflow (@nextflow-io).

Here's how to get started:

- 📚 See available pipelines on the website.
- 🏫 Check out the documentation.
- 💬 Come and say hi on our Slack!
- 🍿 Tune in for talks and events, get involved!
- 🌈 Please abide by our community code of conduct

People

View all

Top languages

Nextflow · Python · Shell · JavaScript · Java

Most used topics

nextflow · nf-core · workflow · pipeline · bioinformatics

Report abuse

nf-core

- nf-core/modules has 500+ ready to use DSL2 modules
  - Rapid prototyping and pipeline building
  - Version controlled, extensive logging, supports Conda, Docker and Singularity

Nextflow DSL2 Module

nf-core/modules

Browse the 1257 modules that are currently available as part of nf-core.

Search

Name

abacas

contiguate draft genome assembly

genome assembly contiguate

Included in: viralrecon

abricate_run

Screen assemblies for antimicrobial resistance against multiple databases

bacteria assembly antimicrobial resistance

Included in: funcscan

abricate_summary

Screen assemblies for antimicrobial resistance against multiple databases

bacteria assembly antimicrobial reistance

abritamr_run

A NATA accredited tool for reporting the presence of antimicrobial resistance genes in bacterial genomes

bacteria fasta antibiotic resistance

adapterremoval

Trim sequencing adapters and collapse overlapping reads

trimming adapters merging fastq

Included in: mag taxprofiler

adapterremovalfixprefix

Fixes prefixes from AdapterRemoval2 output to make sure no clashing read names are in the output. For use with DeDup.

adapterremoval fastq dedup

admixture

ADMIXTURE is a program for estimating ancestry in a model-based manner from large autosomal SNP genotype datasets, where the individuals are unrelated (for example, the individuals in a case-control association study).

ancestry population genetics admixture reference panels gwas

affy_justrma

Read CEL files into an ExpressionSet and generate a matrix

affy microarray expression matrix

Included in: differentialabundance

nf-core

A community effort to collect a curated set of analysis pipelines built using Nextflow.

Home Pipelines Resources Docs Community About Search Join nf-core

callingcards ✓ ☆ 2

A pipeline for processing calling cards data

1.0.0 released 3 days ago

fastquorum ✓ ☆ 11

Pipeline to produce consensus reads using unique molecular indexes/barcodes (UMIs)

consensus umi umis unique-molecular-identifier

1.0.0 released 4 days ago

epitopeprediction ✓ ☆ 37

A bioinformatics best-practice analysis pipeline for epitope prediction and annotation

epitope epitope-prediction mhc-binding-prediction

2.3.1 released 10 days ago

mag ✓ ☆ 182

Assembly and binning of metagenomes

annotation assembly binning long-read-sequencing metagenomes metagenomics nanopore nanopore-sequencing

3.0.0 released 14 days ago

rnasplice ✓ ☆ 34

rnasplice is a bioinformatics pipeline for RNA-seq alternative splicing analysis

alternative-splicing rna rna-seq splicing

1.0.4 released 18 days ago

differentialabundance ✓ ☆ 46

Differential abundance analysis for feature/observation matrices from platforms such as RNA-seq

atac-seq chip-seq deseq2 differential-abundance differential-expression gsea limma microarray rna-seq shiny

1.5.0 released 19 days ago

scrnaseq ✓ ☆ 172

A single-cell RNAseq pipeline for 10X genomics data

10x-genomics 10xgen cellranger kallisto star-solo

2.6.0 released 20 days ago

sarek ✓ ☆ 341

Analysis pipeline to detect germline or somatic variants (pre-processing, variant calling and annotation) from WGS / targeted sequencing

whole-genome-sequencing

3.4.2 released 20 days ago

Over 90 actively maintained workflows
Over 60 workflows with stable releases
https://nf-co.re/pipelines

6. Awesome Workflows

BACTOPIA

COMPLETE ANALYSIS OF BACTERIAL GENOMES

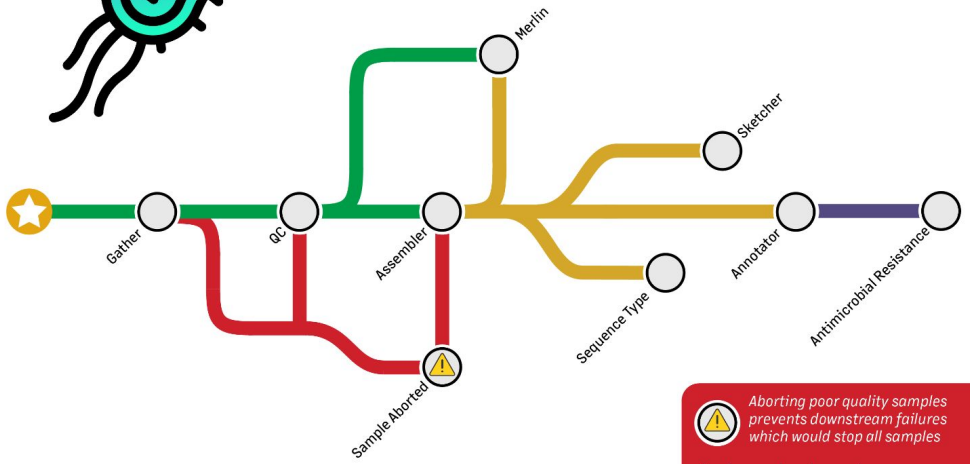

Accepted Inputs

Illumina and/or Nanopore Reads
--R1/--R2, --SE, --ont. --hybrid, --sample
--samples with 'bactopia prepare' file-of-filenames

DDBJ/ENA/SRA Accessions
--accession 'Experiment Accession'
--accessions with 'bactopia search' results

Assemblies
--assembly, --sample
--samples with 'bactopia prepare' file-of-filenames

NCBI Assembly Accessions
--accession 'Assembly Accession'
--accessions 'file with accessions'

Aborting poor quality samples prevents downstream failures which would stop all samples

Too few reads or basepairs

Coverage below minimum

Paired-end with different read counts, mismatched IDs, or skewed proportions

Genome size outside expectation

0 assembled contigs

Assembled size below minimum

Legend

- Process uses FASTQs
- Process uses Contigs
- Process uses Contigs and Proteins
- Minimum QC not met, sample aborted

Bactopia Processes

Gather
Collect local files and/or download sequences from SRA/ENA/DDBJ or NCBI Assembly accessions

QC
Trim and filter low quality reads, subsample to specified coverage, and generate summary metrics

Assembler
Create a de novo assembly (standard, hydbrid, or short read polished) and summary metrics

Merlin
Use Mash distances to automatically execute species specific tools, requires --ask_merlin

Sketcher
Create minmer sketches and query them against GTDB and RefSeq

Annotator
Predict genes and proteins from the assembled contigs

Antimicrobial Resistance
Identify presence of AMR and/or virulence genes

https://github.com/bactopia/bactopia

README | Code of conduct | Apache-2.0 license

🔥🦠🔥 PHoeNIx: A short-read pipeline for healthcare-associated and antimicrobial resistant pathogens

DOI 10.5281/zenodo.10869898

nextflow DSL2 ≥21.10.3 | run with docker | run with singularity

slack StaPH-B #phoenix-dev

Metrics

Page Hits 24 / 14398

🔥🦠🔥 GitHub Clones: 262

For full documentation on the pipeline see the Wiki, but quick start instructions are provided below if you are feeling brave.

https://github.com/CDCgov/phoenix

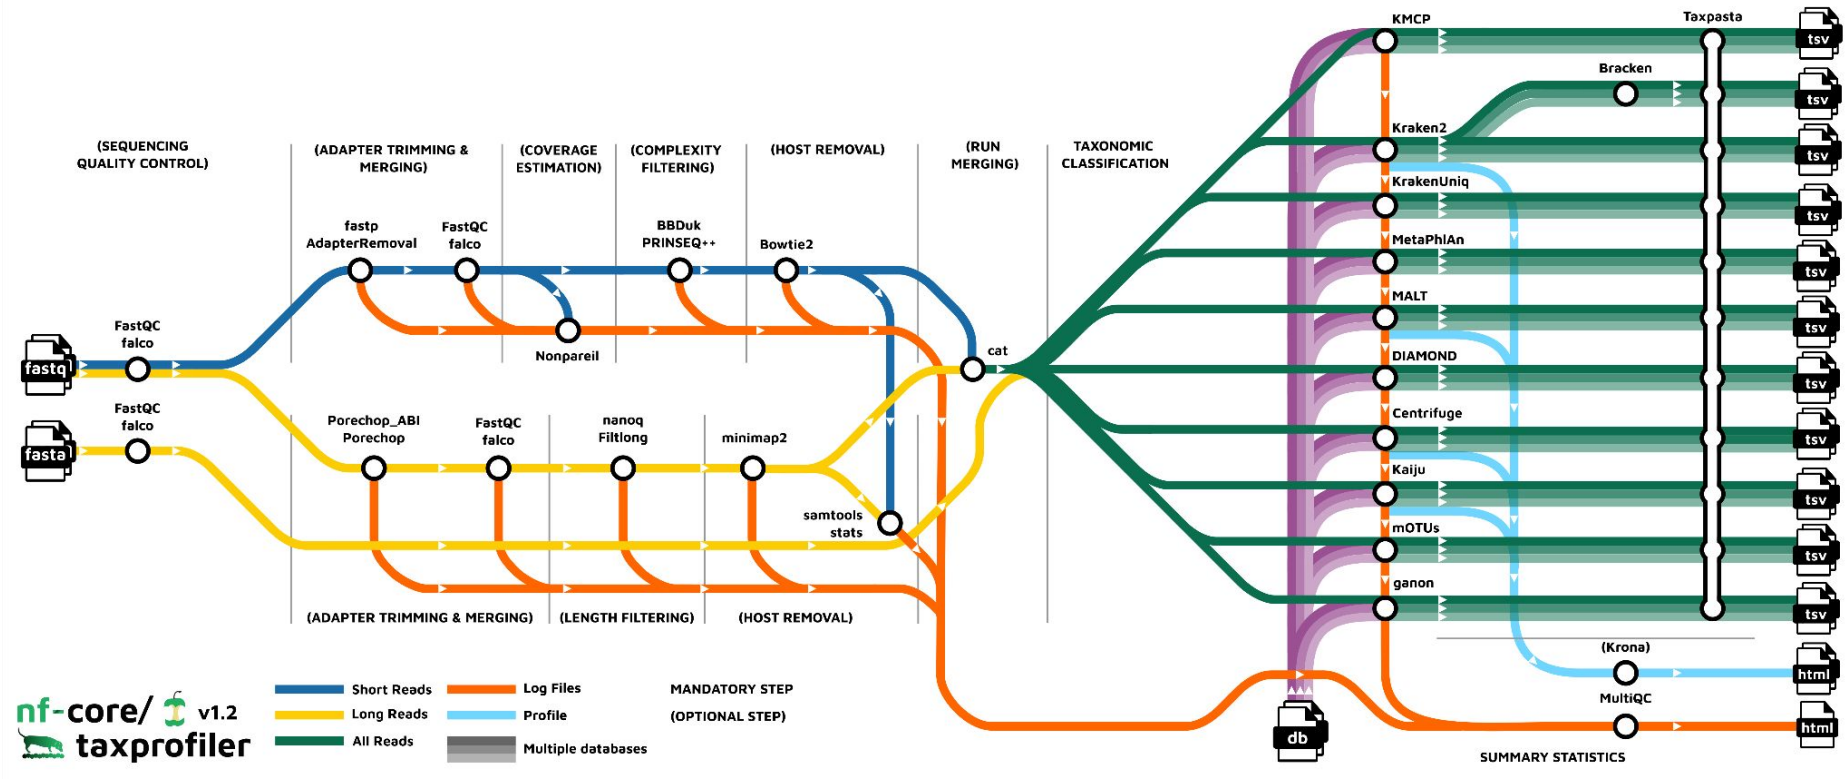

https://github.com/nf-core/taxprofiler

README | Code of conduct | Apache-2.0 license

nf-core/ mycosnp

nextflow DSL2 ≥21.10.3 | run with conda | run with docker | run wi

https://github.com/CDCgov/mycosnp-nf

Hands-On Exercise

Theiagen®

Exercise 04: Writing Your First Nextflow Workflow

Exercise Goal

1. **Use GitHub and your dev environment to:**
   a. Install miniconda
   b. Install and test Nextflow
   c. Wrap fastq-peek.sh into a Nextflow workflow

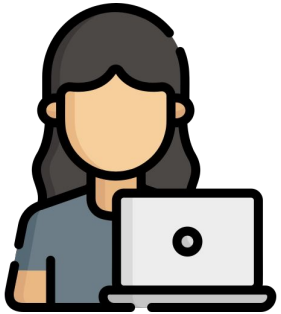

Theiagen